EN3160 Assignment 3 - Neural Networks

Name: Perera U. I. Himal

Index No: 230486A

GitHub: https://github.com/Isira-Himal-Perera/EN3160-Image-Processing-and-Computer-Vision/tree/main/assignment-03

Question 1: Dense Neural Network with Manual Backpropagation

In [13]:
```
import torch
import torchvision
import torchvision.transforms as transforms
import matplotlib.pyplot as plt

# Set random seed for reproducibility
torch.manual_seed(42)

# Select device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Using:", device)

# Load CIFAR-10 dataset
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

trainset = torchvision.datasets.CIFAR10(
    root="./data", train=True, download=True, transform=transform
)
testset = torchvision.datasets.CIFAR10(
    root="./data", train=False, download=True, transform=transform
)

trainloader = torch.utils.data.DataLoader(
    trainset, batch_size=50, shuffle=True,
    num_workers=2, pin_memory=(device.type == "cuda"),
    persistent_workers=True
)
testloader = torch.utils.data.DataLoader(
    testset, batch_size=50, shuffle=False,
    num_workers=2, pin_memory=(device.type == "cuda"),
    persistent_workers=True
)

# Initialize network parameters on the selected device
Din = 3 * 32 * 32  # Input size: 3 channels * 32 pixels * 32 pixels = 3072
H = 100            # Hidden layer nodes (Requirement 1a)
```

```
K = 10              # Output classes for CIFAR-10

w1 = torch.randn(Din, H, device=device) * 0.01
b1 = torch.zeros(H, device=device)

w2 = torch.randn(H, K, device=device) * 0.01
b2 = torch.zeros(K, device=device)

epochs = 10         # Requirement 1c
lr = 0.4
lr_decay = 0.9

loss_history = []

# Train using manual forward propagation and backpropagation (Requirement 1b)
with torch.no_grad():
    for epoch in range(epochs):
        total_loss = torch.zeros((), device=device)

        for images, labels in trainloader:
            images = images.to(device, non_blocking=True)
            labels = labels.to(device, non_blocking=True)

            N = images.size(0)
            x = images.reshape(N, -1)
            indices = torch.arange(N, device=device)

            # Forward propagation
            h = torch.sigmoid(x @ w1 + b1)
            logits = h @ w2 + b2

            # Numerically stable log-softmax
            shifted_logits = logits - logits.max(dim=1, keepdim=True).values
            log_probs = shifted_logits - torch.log(
                torch.exp(shifted_logits).sum(dim=1, keepdim=True)
            )

            # Cross-entropy loss
            loss = -log_probs[indices, labels].mean()
            total_loss += loss * N

            # Backpropagation: output layer
            probs = torch.exp(log_probs)
            dlogits = probs.clone()
            dlogits[indices, labels] -= 1
            dlogits /= N

            dw2 = h.T @ dlogits
            db2 = dlogits.sum(dim=0)

            # Backpropagation: sigmoid hidden layer
            dh = dlogits @ w2.T
            dz1 = dh * h * (1 - h)

            dw1 = x.T @ dz1
            db1 = dz1.sum(dim=0)
```

```
            # Gradient descent update
            w1 -= lr * dw1
            b1 -= lr * db1
            w2 -= lr * dw2
            b2 -= lr * db2

        epoch_loss = total_loss.item() / len(trainset)
        loss_history.append(epoch_loss)

        print(
            f"Epoch {epoch + 1:2d}/{epochs} | "
            f"Training Loss: {epoch_loss:.4f}"
        )

        lr *= lr_decay


# Evaluate accuracy after training (Requirement 1c)
@torch.no_grad()
def accuracy(loader):
    correct = torch.zeros((), dtype=torch.long, device=device)
    total = 0

    for images, labels in loader:
        images = images.to(device, non_blocking=True)
        labels = labels.to(device, non_blocking=True)

        x = images.reshape(images.size(0), -1)

        h = torch.sigmoid(x @ w1 + b1)
        logits = h @ w2 + b2
        predictions = logits.argmax(dim=1)

        correct += (predictions == labels).sum()
        total += labels.size(0)

    return 100 * correct.item() / total


train_acc = accuracy(trainloader)
test_acc = accuracy(testloader)

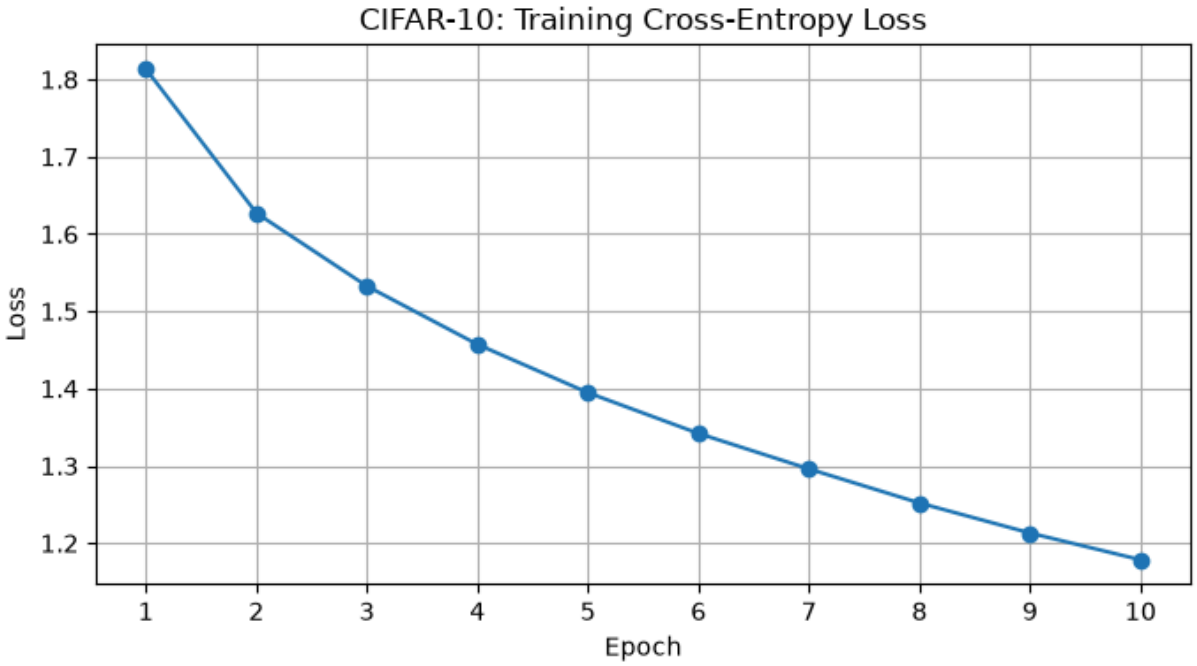
print(f"\nTraining accuracy after 10 epochs: {train_acc:.2f}%")
print(f"Test accuracy after 10 epochs: {test_acc:.2f}%")


# Plot training loss history
plt.figure(figsize=(7, 4))
plt.plot(range(1, epochs + 1), loss_history, "o-")
plt.title("CIFAR-10: Training Cross-Entropy Loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.xticks(range(1, epochs + 1))
plt.grid(True)
```

```
plt.tight_layout()
plt.show()
```

```
Using: cuda
Epoch  1/10 | Training Loss: 1.8135
Epoch  2/10 | Training Loss: 1.6271
Epoch  3/10 | Training Loss: 1.5325
Epoch  4/10 | Training Loss: 1.4572
Epoch  5/10 | Training Loss: 1.3946
Epoch  6/10 | Training Loss: 1.3419
Epoch  7/10 | Training Loss: 1.2960
Epoch  8/10 | Training Loss: 1.2519
Epoch  9/10 | Training Loss: 1.2130
Epoch 10/10 | Training Loss: 1.1784

Training accuracy after 10 epochs: 60.07%
Test accuracy after 10 epochs: 48.42%
```



Results

Following 10 epochs of training, the model demonstrated steady learning progress, with the loss dropping from 1.8135 down to 1.1784. This optimization yielded a training accuracy of 60.07% and a test accuracy of 48.42%.

The relatively modest performance can be attributed to the flattening of the input images, which strips away vital spatial structures that a standard dense layer architecture cannot effectively capture. Furthermore, the performance gap observed between the training and test sets indicates a degree of overfitting.

Question 2: Convolutional Neural Networks (LeNet-5 and Modern CNN) on MNIST

In [25]:

```
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader
from torchvision import datasets, transforms


# Define LeNet-5 Architecture for MNIST
class LeNet5(nn.Module):

  def __init__(self):
    super(LeNet5, self).__init__()
    self.conv1 = nn.Conv2d(1, 6, kernel_size=5)
    self.relu1 = nn.ReLU()
    self.pool1 = nn.AvgPool2d(kernel_size=2, stride=2)

    self.conv2 = nn.Conv2d(6, 16, kernel_size=5)
    self.relu2 = nn.ReLU()
    self.pool2 = nn.AvgPool2d(kernel_size=2, stride=2)

    self.fc1 = nn.Linear(16 * 4 * 4, 120)
    self.relu3 = nn.ReLU()
    self.fc2 = nn.Linear(120, 84)
    self.relu4 = nn.ReLU()
    self.fc3 = nn.Linear(84, 10)

  def forward(self, x):
    x = self.pool1(self.relu1(self.conv1(x)))
    x = self.pool2(self.relu2(self.conv2(x)))
    x = x.view(-1, 16 * 4 * 4)
    x = self.relu3(self.fc1(x))
    x = self.relu4(self.fc2(x))
    x = self.fc3(x)
    return x


device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

transform = transforms.Compose(
    [transforms.ToTensor(), transforms.Normalize((0.1307,), (0.3081,))]
)

train_dataset = datasets.MNIST(
    root="./data", train=True, download=True, transform=transform
)
test_dataset = datasets.MNIST(
    root="./data", train=False, download=True, transform=transform
)

train_loader = DataLoader(dataset=train_dataset, batch_size=64, shuffle=True)
test_loader = DataLoader(dataset=test_dataset, batch_size=1000, shuffle=False)

model = LeNet5().to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)
```

```python
def evaluate_accuracy(loader):
  model.eval()
  correct = 0
  total = 0
  with torch.no_grad():
    for images, labels in loader:
      images, labels = images.to(device), labels.to(device)
      outputs = model(images)
      _, predicted = torch.max(outputs.data, 1)
      total += labels.size(0)
      correct += (predicted == labels).sum().item()
  return 100.0 * correct / total


# Track history for plotting
lenet_train_accs, lenet_test_accs, lenet_losses = [], [], []

num_epochs = 10
print(f"Training LeNet-5 on device: {device}\n")

for epoch in range(num_epochs):
  model.train()
  running_loss = 0.0
  for images, labels in train_loader:
    images, labels = images.to(device), labels.to(device)
    optimizer.zero_grad()
    outputs = model(images)
    loss = criterion(outputs, labels)
    loss.backward()
    optimizer.step()
    running_loss += loss.item()

  train_acc = evaluate_accuracy(train_loader)
  test_acc = evaluate_accuracy(test_loader)

  # Save metrics
  lenet_train_accs.append(train_acc)
  lenet_test_accs.append(test_acc)
  lenet_losses.append(running_loss / len(train_loader))

  print(
      f"Epoch [{epoch+1}/{num_epochs}] | Loss:"
      f" {running_loss/len(train_loader):.4f} | Train Accuracy:"
      f" {train_acc:.2f}% | Test Accuracy: {test_acc:.2f}%"
  )

print("\nLeNet-5 Training complete!")
```

```
Training LeNet-5 on device: cuda

Epoch [1/10] | Loss: 0.3162 | Train Accuracy: 96.76% | Test Accuracy: 96.74%
Epoch [2/10] | Loss: 0.0912 | Train Accuracy: 97.83% | Test Accuracy: 97.78%
Epoch [3/10] | Loss: 0.0617 | Train Accuracy: 98.68% | Test Accuracy: 98.55%
Epoch [4/10] | Loss: 0.0486 | Train Accuracy: 98.52% | Test Accuracy: 98.58%
Epoch [5/10] | Loss: 0.0393 | Train Accuracy: 98.92% | Test Accuracy: 98.61%
Epoch [6/10] | Loss: 0.0353 | Train Accuracy: 99.15% | Test Accuracy: 98.88%
Epoch [7/10] | Loss: 0.0291 | Train Accuracy: 99.38% | Test Accuracy: 99.01%
Epoch [8/10] | Loss: 0.0253 | Train Accuracy: 99.40% | Test Accuracy: 98.98%
Epoch [9/10] | Loss: 0.0236 | Train Accuracy: 99.48% | Test Accuracy: 98.96%
Epoch [10/10] | Loss: 0.0199 | Train Accuracy: 99.52% | Test Accuracy: 98.92%

LeNet-5 Training complete!
```

Discussion: Analysis of Overfitting After Epoch 7

Observing the training progression, the model achieves its peak generalization performance at **Epoch 7**, yielding a maximum test accuracy of **99.01%** alongside a training accuracy of **99.38%**.

Beyond Epoch 7, clear indications of mild overfitting begin to emerge:

- **Divergence in Metrics:** From Epoch 8 through Epoch 10, the training accuracy continues to climb, reaching up to **99.52%** by Epoch 10 while the test accuracy fluctuates and drops slightly below its peak, settling at **98.92%** by Epoch 10.
- **Continuous Loss Reduction:** The training loss continues to decrease steadily down to **0.0199** by the final epoch. This demonstrates that the network is continuing to fit the training distribution closer, but this extra optimization fails to translate into better performance on unseen test data.

Conclusion: To prevent this minor generalization drop, early stopping around **Epoch 7** would help maintain optimal test performance.

In [26]:
```python
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader
from torchvision import datasets, transforms


class ModernCNN(nn.Module):

  def __init__(self):
    super(ModernCNN, self).__init__()
    self.conv1 = nn.Conv2d(1, 32, kernel_size=3, padding=1)
    self.bn1 = nn.BatchNorm2d(32)
    self.relu1 = nn.ReLU()
    self.pool1 = nn.MaxPool2d(kernel_size=2, stride=2)

    self.conv2 = nn.Conv2d(32, 64, kernel_size=3, padding=1)
    self.bn2 = nn.BatchNorm2d(64)
```

```
    self.relu2 = nn.ReLU()
    self.pool2 = nn.MaxPool2d(kernel_size=2, stride=2)

    self.conv3 = nn.Conv2d(64, 128, kernel_size=3, padding=1)
    self.bn3 = nn.BatchNorm2d(128)
    self.relu3 = nn.ReLU()
    self.dropout3 = nn.Dropout2d(p=0.25)

    self.fc1 = nn.Linear(128 * 7 * 7, 128)
    self.relu4 = nn.ReLU()
    self.dropout4 = nn.Dropout(p=0.5)
    self.fc2 = nn.Linear(128, 10)

  def forward(self, x):
    x = self.pool1(self.relu1(self.bn1(self.conv1(x))))
    x = self.pool2(self.relu2(self.bn2(self.conv2(x))))
    x = self.dropout3(self.relu3(self.bn3(self.conv3(x))))
    x = x.view(-1, 128 * 7 * 7)
    x = self.relu4(self.fc1(x))
    x = self.dropout4(x)
    x = self.fc2(x)
    return x


model2 = ModernCNN().to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model2.parameters(), lr=0.001)


def evaluate_accuracy_model2(loader):
  model2.eval()
  correct = 0
  total = 0
  with torch.no_grad():
    for images, labels in loader:
      images, labels = images.to(device), labels.to(device)
      outputs = model2(images)
      _, predicted = torch.max(outputs.data, 1)
      total += labels.size(0)
      correct += (predicted == labels).sum().item()
  return 100.0 * correct / total


# Track history for plotting
modern_train_accs, modern_test_accs, modern_losses = [], [], []

print(f"Training Modern CNN on device: {device}\n")

for epoch in range(num_epochs):
  model2.train()
  running_loss = 0.0
  for images, labels in train_loader:
    images, labels = images.to(device), labels.to(device)
    optimizer.zero_grad()
    outputs = model2(images)
    loss = criterion(outputs, labels)
```

```
    loss.backward()
    optimizer.step()
    running_loss += loss.item()

  train_acc = evaluate_accuracy_model2(train_loader)
  test_acc = evaluate_accuracy_model2(test_loader)

  # Save metrics
  modern_train_accs.append(train_acc)
  modern_test_accs.append(test_acc)
  modern_losses.append(running_loss / len(train_loader))

  print(
      f"Epoch [{epoch+1}/{num_epochs}] | Loss:"
      f" {running_loss/len(train_loader):.4f} | Train Accuracy:"
      f" {train_acc:.2f}% | Test Accuracy: {test_acc:.2f}%"
  )

print("\nModern CNN Training complete!")
```

```
Training Modern CNN on device: cuda

Epoch [1/10] | Loss: 0.1601 | Train Accuracy: 97.19% | Test Accuracy: 97.14%
Epoch [2/10] | Loss: 0.0813 | Train Accuracy: 98.89% | Test Accuracy: 98.80%
Epoch [3/10] | Loss: 0.0625 | Train Accuracy: 99.31% | Test Accuracy: 99.03%
Epoch [4/10] | Loss: 0.0501 | Train Accuracy: 99.47% | Test Accuracy: 99.21%
Epoch [5/10] | Loss: 0.0450 | Train Accuracy: 99.49% | Test Accuracy: 99.26%
Epoch [6/10] | Loss: 0.0388 | Train Accuracy: 99.56% | Test Accuracy: 99.17%
Epoch [7/10] | Loss: 0.0328 | Train Accuracy: 99.56% | Test Accuracy: 99.19%
Epoch [8/10] | Loss: 0.0290 | Train Accuracy: 99.69% | Test Accuracy: 99.26%
Epoch [9/10] | Loss: 0.0263 | Train Accuracy: 99.73% | Test Accuracy: 99.29%
Epoch [10/10] | Loss: 0.0236 | Train Accuracy: 99.76% | Test Accuracy: 99.28%

Modern CNN Training complete!
```

Discussion: Analysis of Overfitting and Performance (Modern CNN)

Observing the training progression, the Modern CNN architecture achieves its optimal test performance around **Epoch 9**, reaching a peak test accuracy of **99.29%** alongside a training accuracy of **99.73%**.

Throughout the 10 epochs, the following dynamics are observed:

- **Rapid Initial Convergence:** The model converges very quickly, exceeding **99.00%** test accuracy as early as **Epoch 3** (99.03%) with a low training loss of **0.0625**.
- **Plateau in Generalization:** From Epoch 4 onwards, test accuracy plateaus in a narrow range between **99.17%** and **99.29%**, while training accuracy continues to climb gradually—reaching **99.76%** by Epoch 10.
- **Consistent Loss Decay:** Training loss steadily decreases throughout training, dropping to **0.0236** by the final epoch, indicating strong optimization without severe overfitting.

Conclusion: The Modern CNN exhibits excellent feature learning with minimal variance between train and test metrics. Implementing early stopping around **Epoch 5 to 7** (where test accuracy remains near **99.26%**) would preserve virtually identical generalization while reducing unnecessary training cycles.

In [27]:
```
import matplotlib.pyplot as plt

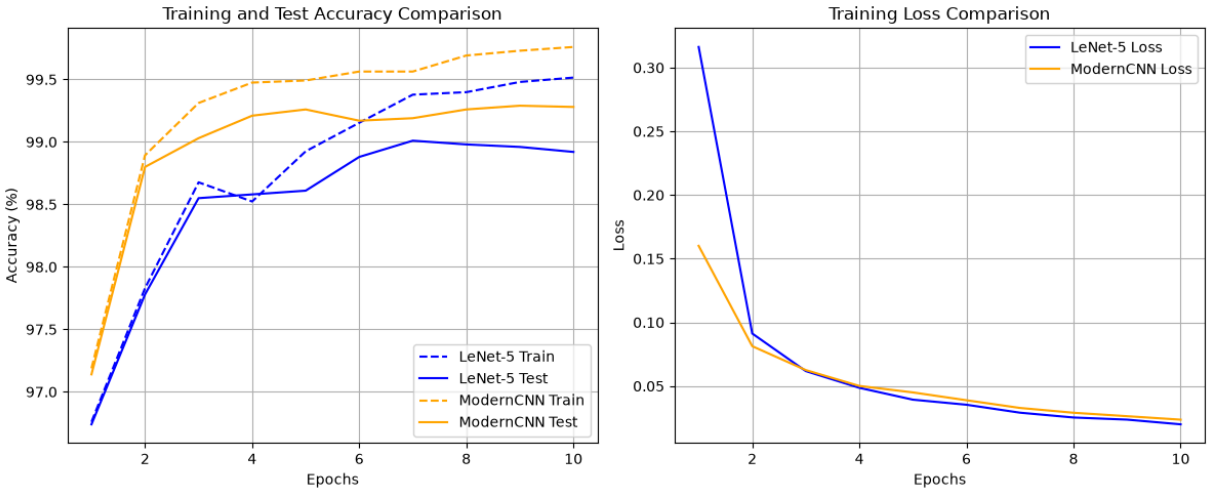
plt.figure(figsize=(12, 5))

# Plot Accuracy Curves
plt.subplot(1, 2, 1)
plt.plot(
    range(1, num_epochs + 1),
    lenet_train_accs,
    label="LeNet-5 Train",
    linestyle="--",
    color="blue",
)
plt.plot(
    range(1, num_epochs + 1),
    lenet_test_accs,
    label="LeNet-5 Test",
    color="blue",
)
plt.plot(
    range(1, num_epochs + 1),
    modern_train_accs,
    label="ModernCNN Train",
    linestyle="--",
    color="orange",
)
plt.plot(
    range(1, num_epochs + 1),
    modern_test_accs,
    label="ModernCNN Test",
    color="orange",
)
plt.xlabel("Epochs")
plt.ylabel("Accuracy (%)")
plt.title("Training and Test Accuracy Comparison")
plt.legend()
plt.grid(True)

# Plot Loss Curves
plt.subplot(1, 2, 2)
plt.plot(
    range(1, num_epochs + 1),
    lenet_losses,
    label="LeNet-5 Loss",
    color="blue",
)
plt.plot(
    range(1, num_epochs + 1),
    modern_losses,
```

```
    label="ModernCNN Loss",
    color="orange",
)
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title("Training Loss Comparison")
plt.legend()
plt.grid(True)

plt.tight_layout()
plt.show()
```



Question 3: Transfer Learning with ResNet-18 on the Hymenoptera Dataset

In [28]:

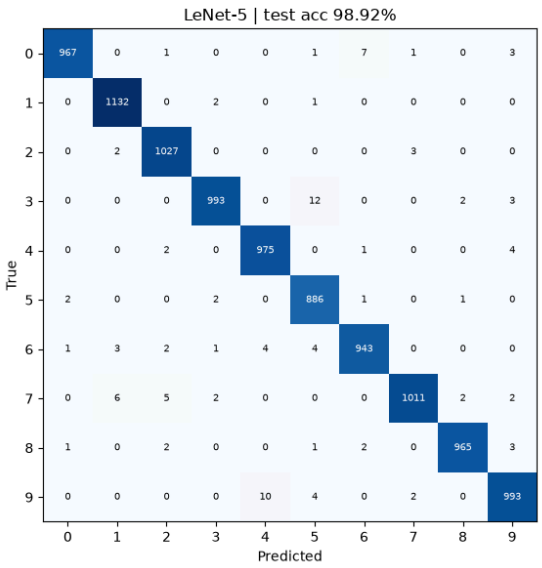
```
import numpy as np

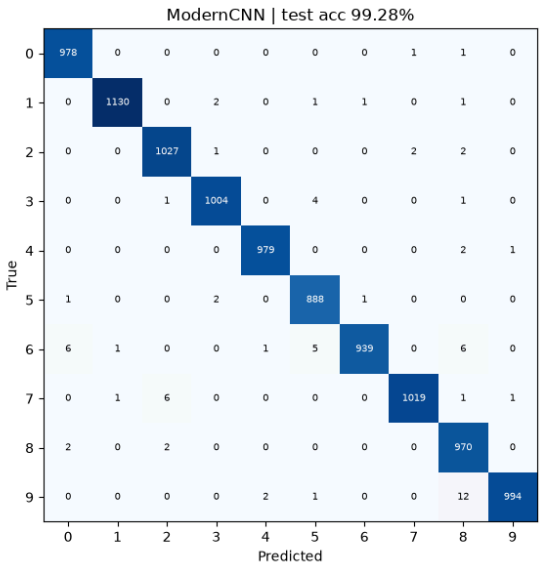

@torch.no_grad()
def predict(net, loader):
  net.eval()
  preds, labels_all = [], []
  for data, labels in loader:
    data = data.to(device)
    preds.append(net(data).argmax(1).cpu())
    labels_all.append(labels)
  return torch.cat(preds).numpy(), torch.cat(labels_all).numpy()


def conf_matrix(preds, labels, n=10):
  cm = np.zeros((n, n), dtype=int)
  np.add.at(cm, (labels, preds), 1)  # rows = true class, cols = predicted class
  return cm


fig, ax = plt.subplots(1, 2, figsize=(14, 6))
for a, (name, net) in zip(ax, [("LeNet-5", model), ("ModernCNN", model2)]):
  p, l = predict(net, test_loader)
  cm = conf_matrix(p, l)
```

```
  a.imshow(cm, cmap="Blues")
  a.set(
      title=f"{name} | test acc {100*(p==l).mean():.2f}%",
      xlabel="Predicted",
      ylabel="True",
      xticks=range(10),
      yticks=range(10),
  )
  for i in range(10):
    for j in range(10):
      a.text(
          j,
          i,
          cm[i, j],
          ha="center",
          va="center",
          fontsize=7,
          color="white" if cm[i, j] > cm.max() / 2 else "black",
      )
plt.tight_layout()
plt.show()
```





In [33]:

```
import os
import urllib.request
import zipfile
import matplotlib.pyplot as plt
import numpy as np
import torch
import torch.nn as nn
import torchvision.transforms as transforms
from torchvision import datasets, models

EPOCHS = 15
BATCH_SIZE = 16
LEARNING_RATE = 0.001
MOMENTUM = 0.9
STEP_SIZE = 7
GAMMA = 0.1
```

```
SEEDS = [0, 1, 2]

DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
CRITERION = nn.CrossEntropyLoss()

DATA_DIR = "hymenoptera_data"
URL = "https://download.pytorch.org/tutorial/hymenoptera_data.zip"

if not os.path.isdir(DATA_DIR):
    urllib.request.urlretrieve(URL, "hymenoptera_data.zip")
    with zipfile.ZipFile("hymenoptera_data.zip", "r") as z:
        z.extractall(".")

MEAN, STD = (0.485, 0.456, 0.406), (0.229, 0.224, 0.225)

train_transforms = transforms.Compose([
    transforms.RandomResizedCrop(224),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize(MEAN, STD),
])

val_transforms = transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(MEAN, STD),
])

train_dataset = datasets.ImageFolder(os.path.join(DATA_DIR, "train"), train_transfo
train_dataset_eval = datasets.ImageFolder(os.path.join(DATA_DIR, "train"), val_tran
val_dataset = datasets.ImageFolder(os.path.join(DATA_DIR, "val"), val_transforms)

class_names = train_dataset.classes
num_classes = len(class_names)

train_loader = torch.utils.data.DataLoader(train_dataset, batch_size=BATCH_SIZE, sh
train_loader_eval = torch.utils.data.DataLoader(train_dataset_eval, batch_size=BATC
val_loader = torch.utils.data.DataLoader(val_dataset, batch_size=BATCH_SIZE, shuffl

def build_transfer_model(mode="finetune"):
    model = models.resnet18(weights=models.ResNet18_Weights.IMAGENET1K_V1)
    if mode == "extractor":
        for param in model.parameters():
            param.requires_grad = False  # Freeze backbone features
    model.fc = nn.Linear(model.fc.in_features, num_classes)
    return model.to(DEVICE)

@torch.no_grad()
def evaluate_model(model, loader, return_predictions=False):
    model.eval()
    total_loss, correct_preds, total_samples = 0.0, 0, 0
    preds_list, labels_list = [], []

    for inputs, targets in loader:
        inputs, targets = inputs.to(DEVICE), targets.to(DEVICE)
```

```
        outputs = model(inputs)
        total_loss += CRITERION(outputs, targets).item() * targets.size(0)
        preds = outputs.argmax(dim=1)
        correct_preds += (preds == targets).sum().item()
        total_samples += targets.size(0)
        if return_predictions:
            preds_list.append(preds.cpu())
            labels_list.append(targets.cpu())

    avg_loss = total_loss / total_samples
    accuracy = 100.0 * correct_preds / total_samples
    if return_predictions:
        return avg_loss, accuracy, torch.cat(preds_list).numpy(), torch.cat(labels_
    return avg_loss, accuracy

def train_model(model, mode, verbose=True):
    params = [p for p in model.parameters() if p.requires_grad]
    optimizer = torch.optim.SGD(params, lr=LEARNING_RATE, momentum=MOMENTUM)
    scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=STEP_SIZE, gam

    history = {"train_loss": [], "val_loss": [], "train_acc": [], "val_acc": []}
    for epoch in range(EPOCHS):
        model.eval() if mode == "extractor" else model.train()
        for inputs, targets in train_loader:
            inputs, targets = inputs.to(DEVICE), targets.to(DEVICE)
            optimizer.zero_grad()
            loss = CRITERION(model(inputs), targets)
            loss.backward()
            optimizer.step()
        scheduler.step()

        train_loss, train_acc = evaluate_model(model, train_loader_eval)
        val_loss, val_acc = evaluate_model(model, val_loader)
        for k, v in zip(history, [train_loss, val_loss, train_acc, val_acc]):
            history[k].append(v)
        if verbose:
            print(f"Epoch {epoch+1:2d}/{EPOCHS} | Train Loss: {train_loss:.4f} | Va
    return history

def compute_confusion_matrix(preds, targets):
    cm = np.zeros((num_classes, num_classes), dtype=int)
    np.add.at(cm, (targets, preds), 1)
    return cm

modes = ["finetune", "extractor"]
experiment_history = {m: [] for m in modes}
final_val_accuracies = {m: [] for m in modes}
final_train_accuracies = {m: [] for m in modes}
confusion_matrices = {}

for mode in modes:
    for seed in SEEDS:
        torch.manual_seed(seed)
        np.random.seed(seed)
        net = build_transfer_model(mode)
        hist = train_model(net, mode, verbose=(seed == SEEDS[0]))
```

```
        experiment_history[mode].append(hist)
        _, val_acc, preds, labels = evaluate_model(net, val_loader, return_predicti
        final_val_accuracies[mode].append(val_acc)
        final_train_accuracies[mode].append(hist["train_acc"][-1])
        if seed == SEEDS[0]:
            confusion_matrices[mode] = compute_confusion_matrix(preds, labels)
        print(f"{mode} | seed {seed} | train acc: {hist['train_acc'][-1]:.2f}% | va
        del net

print("\nFinal Results Summary (Mean ± Std):")
for mode in modes:
    print(f"{mode:10s} | train: {np.mean(final_train_accuracies[mode]):.2f} ± {np.s

colors = {"finetune": "tab:blue", "extractor": "tab:orange"}
ep = range(1, EPOCHS + 1)
fig, ax = plt.subplots(1, 2, figsize=(13, 4.5))

for mode in modes:
    h, c = experiment_history[mode][0], colors[mode]
    ax[0].plot(ep, h["train_loss"], "-o", color=c, markersize=3, label=f"{mode} tra
    ax[0].plot(ep, h["val_loss"], "--s", color=c, markersize=3, label=f"{mode} val"
    ax[1].plot(ep, h["train_acc"], "-o", color=c, markersize=3, label=f"{mode} trai
    ax[1].plot(ep, h["val_acc"], "--s", color=c, markersize=3, label=f"{mode} val")

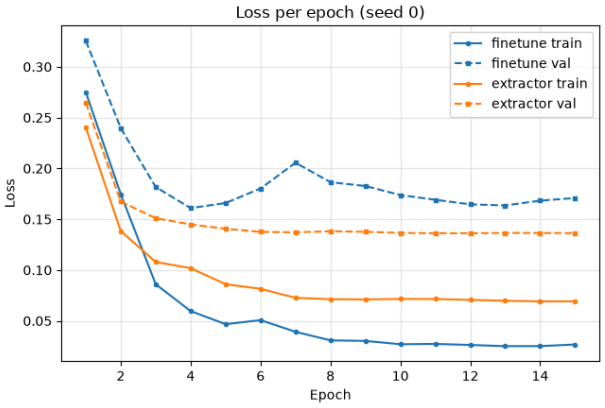
ax[0].set(title="Loss per epoch (seed 0)", xlabel="Epoch", ylabel="Loss")
ax[1].set(title="Accuracy per epoch (seed 0)", xlabel="Epoch", ylabel="Accuracy (%)
for a in ax:
    a.grid(alpha=0.3)
    a.legend()
plt.tight_layout()
plt.show()

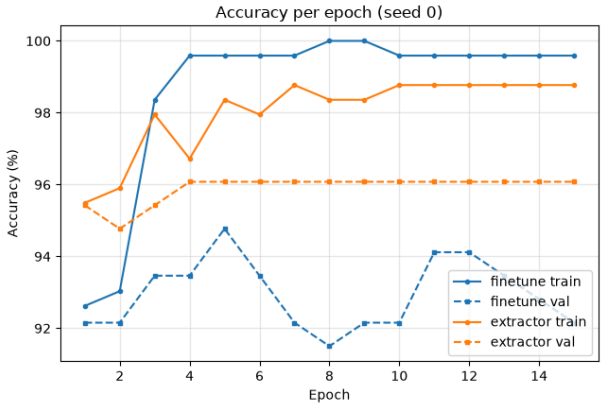
fig, axes = plt.subplots(1, len(confusion_matrices), figsize=(5.5 * len(confusion_m
for a, (mode, cm) in zip(axes[0], confusion_matrices.items()):
    a.imshow(cm, cmap="Blues")
    a.set(title=f"{mode} confusion matrix (seed 0)", xlabel="Predicted", ylabel="Tr
          xticks=range(num_classes), yticks=range(num_classes), xticklabels=class_n
    for i in range(num_classes):
        for j in range(num_classes):
            a.text(j, i, cm[i, j], ha="center", va="center", fontsize=12,
                   color="white" if cm[i, j] > cm.max() / 2 else "black")
plt.tight_layout()
plt.show()
```

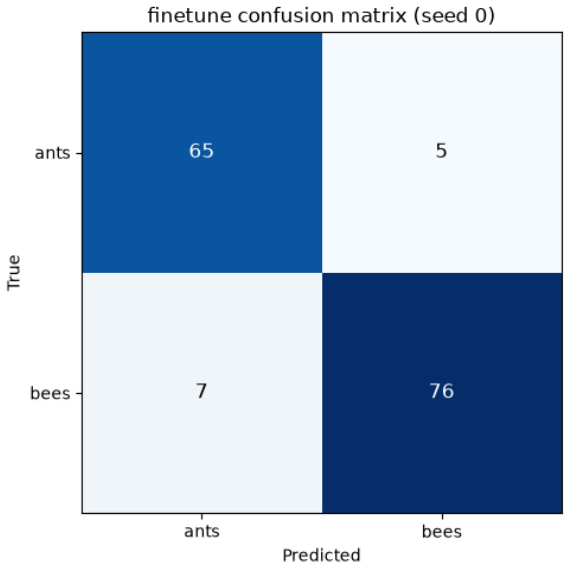
```
Epoch  1/15 | Train Loss: 0.2750 | Val Loss: 0.3255 | Train Acc: 92.62% | Val Acc: 9
2.16%
Epoch  2/15 | Train Loss: 0.1747 | Val Loss: 0.2396 | Train Acc: 93.03% | Val Acc: 9
2.16%
Epoch  3/15 | Train Loss: 0.0860 | Val Loss: 0.1816 | Train Acc: 98.36% | Val Acc: 9
3.46%
Epoch  4/15 | Train Loss: 0.0596 | Val Loss: 0.1609 | Train Acc: 99.59% | Val Acc: 9
3.46%
Epoch  5/15 | Train Loss: 0.0467 | Val Loss: 0.1660 | Train Acc: 99.59% | Val Acc: 9
4.77%
Epoch  6/15 | Train Loss: 0.0507 | Val Loss: 0.1802 | Train Acc: 99.59% | Val Acc: 9
3.46%
Epoch  7/15 | Train Loss: 0.0392 | Val Loss: 0.2056 | Train Acc: 99.59% | Val Acc: 9
2.16%
Epoch  8/15 | Train Loss: 0.0308 | Val Loss: 0.1864 | Train Acc: 100.00% | Val Acc:
91.50%
Epoch  9/15 | Train Loss: 0.0302 | Val Loss: 0.1827 | Train Acc: 100.00% | Val Acc:
92.16%
Epoch 10/15 | Train Loss: 0.0270 | Val Loss: 0.1738 | Train Acc: 99.59% | Val Acc: 9
2.16%
Epoch 11/15 | Train Loss: 0.0273 | Val Loss: 0.1691 | Train Acc: 99.59% | Val Acc: 9
4.12%
Epoch 12/15 | Train Loss: 0.0262 | Val Loss: 0.1646 | Train Acc: 99.59% | Val Acc: 9
4.12%
Epoch 13/15 | Train Loss: 0.0251 | Val Loss: 0.1635 | Train Acc: 99.59% | Val Acc: 9
3.46%
Epoch 14/15 | Train Loss: 0.0251 | Val Loss: 0.1683 | Train Acc: 99.59% | Val Acc: 9
2.81%
Epoch 15/15 | Train Loss: 0.0267 | Val Loss: 0.1709 | Train Acc: 99.59% | Val Acc: 9
2.16%
finetune | seed 0 | train acc: 99.59% | val acc: 92.16%
finetune | seed 1 | train acc: 99.59% | val acc: 94.77%
finetune | seed 2 | train acc: 100.00% | val acc: 94.77%
Epoch  1/15 | Train Loss: 0.2402 | Val Loss: 0.2645 | Train Acc: 95.49% | Val Acc: 9
5.42%
Epoch  2/15 | Train Loss: 0.1383 | Val Loss: 0.1673 | Train Acc: 95.90% | Val Acc: 9
4.77%
Epoch  3/15 | Train Loss: 0.1079 | Val Loss: 0.1510 | Train Acc: 97.95% | Val Acc: 9
5.42%
Epoch  4/15 | Train Loss: 0.1018 | Val Loss: 0.1449 | Train Acc: 96.72% | Val Acc: 9
6.08%
Epoch  5/15 | Train Loss: 0.0860 | Val Loss: 0.1406 | Train Acc: 98.36% | Val Acc: 9
6.08%
Epoch  6/15 | Train Loss: 0.0815 | Val Loss: 0.1376 | Train Acc: 97.95% | Val Acc: 9
6.08%
Epoch  7/15 | Train Loss: 0.0726 | Val Loss: 0.1371 | Train Acc: 98.77% | Val Acc: 9
6.08%
Epoch  8/15 | Train Loss: 0.0713 | Val Loss: 0.1382 | Train Acc: 98.36% | Val Acc: 9
6.08%
Epoch  9/15 | Train Loss: 0.0711 | Val Loss: 0.1376 | Train Acc: 98.36% | Val Acc: 9
6.08%
Epoch 10/15 | Train Loss: 0.0716 | Val Loss: 0.1365 | Train Acc: 98.77% | Val Acc: 9
6.08%
Epoch 11/15 | Train Loss: 0.0715 | Val Loss: 0.1362 | Train Acc: 98.77% | Val Acc: 9
6.08%
Epoch 12/15 | Train Loss: 0.0706 | Val Loss: 0.1362 | Train Acc: 98.77% | Val Acc: 9
```

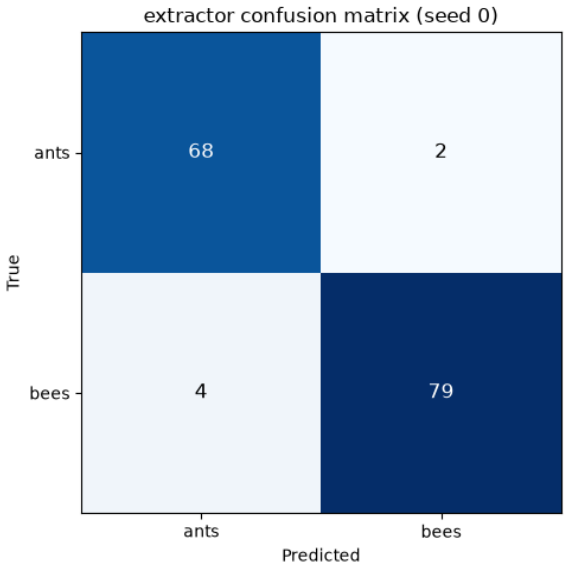
```
6.08%
Epoch 13/15 | Train Loss: 0.0698 | Val Loss: 0.1365 | Train Acc: 98.77% | Val Acc: 9
6.08%
Epoch 14/15 | Train Loss: 0.0691 | Val Loss: 0.1364 | Train Acc: 98.77% | Val Acc: 9
6.08%
Epoch 15/15 | Train Loss: 0.0692 | Val Loss: 0.1363 | Train Acc: 98.77% | Val Acc: 9
6.08%
extractor | seed 0 | train acc: 98.77% | val acc: 96.08%
extractor | seed 1 | train acc: 97.95% | val acc: 96.08%
extractor | seed 2 | train acc: 98.36% | val acc: 96.08%

Final Results Summary (Mean ± Std):
finetune   | train: 99.73 ± 0.19% | val: 93.90 ± 1.23%
extractor  | train: 98.36 ± 0.33% | val: 96.08 ± 0.00%
```









Observations:

1. **Superior Generalization of Feature Extraction:** The feature extractor generalizes better than fine-tuning, achieving a higher and completely stable validation accuracy of **96.08%** across all random seeds with zero variance, compared to fine-tuning's average validation accuracy of **93.90% ± 1.23%**.
2. **Evidence of Overfitting in Fine-Tuning:** Fine-tuning exhibits clear signs of overfitting on the training set, driving training accuracy up to nearly **100%** while causing the

validation loss to drift upward and validation accuracy to fluctuate after the initial epochs.

Conclusion & Discussion:

For a small dataset that shares strong similarities with ImageNet, feature extraction proves to be the superior approach—yielding higher accuracy, perfect stability across seeds, and stronger resistance to overfitting. Fine-tuning would likely require a larger training dataset, a lower learning rate for the backbone layers, or a more significant domain shift away from ImageNet to outperform feature extraction.